

CPT is the more accurate model (%)

— all differing tokens (baseline)

● high-confidence veto (w < 0.1)

